

This is the mail archive of the gcc@gcc.gnu.org mailing list for the [GCC project](https://gcc.gnu.org/).

Index Nav: [\[Date Index\]](#) [\[Subject Index\]](#) [\[Author Index\]](#) [\[Thread Index\]](#)

Message Nav: [\[Date Prev\]](#) [\[Date Next\]](#) [\[Thread Prev\]](#) [\[Thread Next\]](#)

Other format: [\[Raw text\]](#)

A proposal to align GCC stack

- *From:* "Ye, Joey" <joe.ye at intel dot com>
 - *To:* <gcc at gcc dot gnu dot org>
 - *Cc:* "H.J. Lu" <hjl at lucon dot org>, <xuepeng dot guo at intel dot com>
 - *Date:* Tue, 18 Dec 2007 10:25:42 +0800
 - *Subject:* A proposal to align GCC stack
-

-- 0. MOTIVATION --

Some local variables (such as of `__m128` type or marked with alignment attribute) require stack aligned at a boundary larger than the default stack

boundary. Current GCC partially supports this with limitations. We are proposing a new design to fully solve the problem.

-- 1. CURRENT IMPLEMENTATION --

There are two ways current GCC supports bigger than default stack alignment. One is to make sure that stack is aligned at program entry point, and then ensure that for each non-leaf function, its frame size is

aligned. This approach doesn't work when linking with libs or objects compiled by other psABI conforming compilers. Some problems are logged as

PR 33721. Another is to adjust stack alignment at the entry point of a function if it is marked with `__attribute__((force_align_arg_pointer))` or `-mstackrealign` option is provided. This method guarantees the alignment

in most of the cases but with following problems and limitations:

- * Only 16 bytes alignment is supported
- * Adjusting stack alignment at each function prologue hurts performance unnecessarily, because not all functions need bigger alignment. In fact, commonly only those functions which have SSE variables defined locally (either declared by the user or compiler generated internal temporary variables) need corresponding alignment.
- * Doesn't support x86_64 for the cases when required stack alignment is > 16 bytes
- * Emits inefficient and complicated prologue/epilogue code to adjust stack alignment
- * Doesn't work with nested functions
- * Has a bug handling register parameters, which resulted in a cpu2006 failure. A patch is available as a workaround.

-- 2. NEW PROPOSAL: DESIGN --

Here, we propose a new design to fully support stack alignment while overcoming above problems. The new design will

- * Support arbitrary alignment value, including 4, 8, 16, 32...
- * Adjust function stack alignment only when necessary
- * Initial development will be on i386 and x86_64, but can be extended to other platforms
- * Emit more efficient prologue/epilogue code
- * Coexist with special features like dynamic stack allocation (`alloca`),

nested functions, register parameter passing, PIC code and tail call optimization
 * Be able to debug and unwind stack

2.1 Support arbitrary alignment value

Different source code and optimizations requires different stack alignment,
 as in following table:

Feature	Alignment (bytes)
i386_ABI	4
x86_64_ABI	16
char	1
short	2
int	4
long	4/8*
long long	8
__m64	8
__m128	16
float	4
double	8
long double	4/16*
user specified	any power of 2

*Note: 4 for i386, 8/16 for x86_64

The new design will support any alignment value in this table.

2.2 Adjust function stack alignment only when necessary

Current GCC defines following macros related to stack alignment:

- i. STACK_BOUNDARY in bits, which is enforced by hardware, 32 for i386 and 64 for x86_64. It is the minimum stack boundary. It is fixed.
- ii. PREFERRED_STACK_BOUNDARY. It sets the stack alignment when calling a function. It may be set at command line and has no impact on stack alignment at function entry. This proposal requires PREFERRED >= STACK, and by default set to ABI_STACK_BOUNDARY

This design will define a few more macros, or concepts not explicitly defined in code:

- iii. ABI_STACK_BOUNDARY in bits, which is the stack boundary specified by psABI, 32 for i386 and 128 for x86_64. ABI_STACK_BOUNDARY >= STACK_BOUNDARY. It is fixed for a given psABI.
- iv. LOCAL_STACK_BOUNDARY in bits. Each function stack has its own stack alignment requirement, which depends the alignment of its stack variables,
 LOCAL_STACK_BOUNDARY = MAX (alignment of each effective stack variable).
- v. INCOMING_STACK_BOUNDARY in bits, which is the stack boundary at function entry. If a function is marked with `__attribute__((force_align_arg_pointer))` or `-mstackrealign` option is provided, INCOMING = STACK_BOUNDARY. Otherwise,
 INCOMING == MIN(ABI_STACK_BOUNDARY, PREFERRED_STACK_BOUNDARY) because a function can be called via psABI externally or called locally with PREFERRED_STACK_BOUNDARY.
- vi. REQUIRED_STACK_ALIGNMENT in bits, which is stack alignment required by local variables and calling other function. REQUIRED_STACK_ALIGNMENT == MAX(LOCAL_STACK_BOUNDARY, PREFERRED_STACK_BOUNDARY) in case of a non-leaf function. For a leaf function, REQUIRED_STACK_ALIGNMENT == LOCAL_STACK_BOUNDARY.

This proposal won't adjust stack when INCOMING_STACK_BOUNDARY >= REQUIRED_STACK_ALIGNMENT. Only when INCOMING_STACK_BOUNDARY < REQUIRED_STACK_ALIGNMENT, it will adjust stack to REQUIRED_STACK_ALIGNMENT at prologue.

2.3 Initial development on i386 and x86_64

We initially support i386 and x86_64. In this document we focus more on i386 because it is hard to implement because of the restriction of having a small register file. But all that we discuss can be easily applied to x86_64.

2.4 Emit more efficient prologue/epilogue

When a function needs to adjust stack alignment and has no dynamic stack allocation, this design will generate following example

prologue/epilogue

code:

IA32 example Prologue:

```
    pushl    %ebp
    movl     %esp, %ebp
    andl     $-16, %esp
    subl     $4, %esp ; is $-4 the local stack size?
```

Epilogue:

```
    movl     %ebp, %esp
    popl     %ebp
    ret
```

Locals will be addressed as esp + offset and parameters as ebp + offset.

Add x86_64 example here.

Thus BP points to parameter frame and SP points to local frame.

2.5 Coexist with special features

Stack alignment adjustment will coexist with varying GCC features that have special calling conventions and frame layout, such as dynamic stack allocation (alloca), nested functions and parameter passing via registers to local functions.

I386 hard register usage is the major problem to make the proposal friendly to various GCC features. This design requires an additional hard register in prologue/epilogue in case of dynamic stack allocation. Because I386 PIC requires BX as GOT pointer and I386 may use AX, DX and CX as parameter passing registers, there are limited candidates for this proposal to choose. Current proposal suggests EDI, because it won't conflict with i386 PIC or regparm.

X86_64 is much easier. This proposal just chooses RBX.

2.5.1 When stack alignment adjustment comes together with alloca, following

example prologue/epilogue will be emitted:

Prologue:

```
    pushl     %edi                // Save callee save reg edi
    leal      8(%esp), %edi       // Save address of parameter
frame
    andl      $-16, %esp          // Align local stack

// Reserve two stack slots and save return address
// and previous frame pointer into them. By
// pointing new ebp to them, we build a pseudo
// stack for unwinding.
    pushl     $4(%edi)           // save return address
    pushl     %ebp               // save old ebp
    movl      %esp, %ebp         // point ebp to pseudo frame
start

    subl      $24, %esp           // adjust local frame size
    movl      %edi, %regl
```

epilogue:

```

    movl    vregl, %edi
    movl    %ebp, %esp          // Restore esp to pseudo frame
start
    popl    %ebp
    leal    -8(%edi), %esp      // restore esp to real frame
start
    popl    %edi               // Restore edi
    ret

```

Locals will be addressed as `ebp - offset`, parameters as `vregl + offset`

Where BX is used to set up virtual parameter frame pointer, BP points to local frame and SP points to dynamic allocation frame.

2.5.2 Nested functions will automatically work because it uses CX as static pointer, which won't conflict with any registers used by stack alignment adjustment, even when nested functions are called via function pointer and a function stub on stack.

2.5.3 GCC may optimize to use registers to pass parameters . At most AX, DX and CX will be used. Such optimization won't conflict with stack alignment adjustment thus it should automatically work.

2.5.4 I386 PIC uses EBX as GOT pointer. This design work well under i386 PIC:

For example:

i686 Prologue:

```

    pushl    %edi
    leal     8(%esp), %edi
    andl     $-16, %esp
    pushl    $4(%edi)
    pushl    %ebp
    movl     %esp, %ebp
    subl     $24, %esp
    call     .L1
.L1:
    popl     %ebx
    movl     %edi, vregl

```

Body: // code for alloca

```

    movl     (vregl), %eax
    subl     %eax, %esp
    andl     $-16, %esp
    movl     %esp, %eax

```

i686 Epilogue:

```

    movl     %ebp, %esp
    popl     %ebp
    leal     -8(%edi), %esp
    popl     %edi
    ret

```

Locals will be addressed as `ebp - offset`, parameters as `vregl + offset`, ebx has the GOT pointer.

2.6 Debug and unwind will work since DWARF2 has the flexibility to define different frame pointers.

2.7 Some intrinsics rely on stack layout. Need to handle them accordingly.

They are `__builtin_return_address`, `__builtin_frame_address`. This proposal will setup pseudo frame slot to help unwinder find return address and

parent frame address by emit following prologue code after adjusting alignment:

```
    pushl    $4(%edi)
    pushl    %ebp
```

-- 3. NEW PROPOSAL: IMPLEMENTATION --

The proposed implementation can be partitioned into following subtasks.

- * Alignment requirement collection
- * Frames addressing
- * Alignment code generation
- * Debug and unwind information

3.1 Collect alignment requirement

Collecting each function's alignment requirement from frontend or from optimization passes like vectorizer, and informing backend.

Current GCC uses `cfun->stack_alignment_needed` to store `MIN(largest stack variable alignment, PREFERRED_STACK_BOUNDARY)`. We will reuse this field and define its value only as "largest stack variable alignment"

3.2 Frames addressing

Adding parameter frame, local frame, static frame and dynamic frame with appropriate pointers, either hard registers or virtual registers.

Backend will customize `CAN_ELIMINATE` hook to assign hard registers to corresponding virtual registers.

3.3 Alignment code generation

Emit prologue/epilogue code to guarantee correct stack alignment based on each function's alignment requirement collected previously.

Modification should happen in `ix86_expand_prologue` and `ix86_expand_epilogue`.

Code to be emitted can follow above design in a straight forward manner.

3.4 Debug information

Emit debug and unwind information for aligned stacks. It also happens in `ix86_expand_prologue` and `ix86_expand_epilogue` corresponding the prologue/epilogue code emitted.

4. Code Example

Simply function:

```
void foo()
{
    volatile int local;
    ...
}
```

i686 Prologue:

```
    pushl    %ebp
    movl     %esp, %ebp
    subl     $4, %esp          // Adjust local frame size by 4
```

i686 Epilogue:

```
    movl     %ebp, %esp
    popl     %ebp
    ret
```

x86_64 Prologue:

```
    pushq    %rbp
    movq     %rsp, %rbp
    subq     $16, %rsp
```

x86_64 Epilogue:

```

    movl    %rbp, %rsp
    popl    %rbp
    ret

```

Pure 16 bytes align:

```

void foo()
{
    volatile __m128 m = _mm_set_ps1(0.f);
}

```

i686 Prologue:

```

    pushl    %ebp
    movl     %esp, %ebp
    andl     $-16, %esp
    subl     $16, %esp    // this is space for m, 16 byte aligned

```

i686 Epilogue:

```

    movl     %ebp, %esp
    popl     %ebp
    ret

```

x86_64 Prologue:

```

    pushq    %rbp
    movq     %rsp, %rbp
    andq     $-16, %rsp
    subq     $16, %rsp

```

x86_64 Epilogue:

```

    movl     %rbp, %rsp
    popl     %rbp
    ret

```

16 bytes align with alloca:

```

void foo(int size)
{
    char * ptr=alloca(size);
    volatile int __attribute__((aligned(32))) m = 0;
    ...
}

```

i686 Prologue:

```

    pushl    %edi
    leal     8(%esp), %edi
    andl     $-32, %esp
    pushl    $4(%edi)
    pushl    %ebp
    movl     %esp, %ebp
    subl     $24, %esp

```

Body: // code for alloca

```

    movl     %edi, %vreg1
    movl     (%vreg1), %eax
    subl     %eax, %esp
    andl     $-16, %esp
    movl     %esp, %eax

```

i686 Epilogue:

```

    movl     %ebp, %esp
    popl     %ebp
    leal     -8(%edi), %esp
    popl     %edi
    ret

```

```

void foo(int dummy1, int dummy2, int dummy3, int dummy4,
        int dummy5, int dummy6, int size)
{

```

```

    char * ptr=alloca(size);
    volatile int __attribute__((aligned(32))) m = 0;
    ...
}

```

x86_64 Prologue:

```

pushq    %rbx
leaq     $16(%rsp), %rbx
andq     $-32, %rsp
pushq    8(%rbx)
pushq    %rbp
movq     %rsp, %rbp
subq     $24, %rsp

```

Body:

```

movq     %rbx, vreg1
movl     (vreg1), %eax
subq     %rax, %rsp
andq     $-16, %rsp
movq     %rsp, %rax

```

x86_64 Epilogue:

```

movl     %rbp, %rsp
popl     %rbp
movl     %rbx, %rsp
popl     %rbx
ret

```

m128 and PIC

```

int g_i;
void foo()
{
    volatile __m128 m = _mm_set_ps1(0.f);
    g_i = 123;
    ...
}

```

i686 Prologue:

```

pushl    %ebp
movl     %esp, %ebp
andl     $-16, %esp
pushl    %ebx
subl     $16, %esp
call     .L1

```

.L1:

```

popl     %ebx
...

```

i686 Epilogue:

```

addl     $16, %esp
popl     %ebx
movl     %ebp, %esp
popl     %ebp
ret

```

m128 + alloca + PIC

```

void foo(int size)
{
    char * ptr=alloca(size);
    volatile __m128 m = _mm_set_ps1(0.f);
    ...
}

```

i686 Prologue:

```

pushl    %edi
leall    8(%esp), %edi
andl     $-16, %esp
pushl    4(%edi)
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
call     .L1

```

.L1:

```

popl     %ebx

```

Body:

```

    movl    %edi, vreg1
    movl    (vreg1), %eax
    subl    %eax, %esp
    andl    $-16, %esp
    movl    %esp, %eax

```

i686 Epilogue:

```

    movl    %ebp, %esp
    popl    %ebp
    leal    -8(%edi), %esp
    popl    %edi
    ret

```

m128 + alloca + PIC + library call

```

void foo(int size)
{
    char * ptr=alloca(size);
    volatile __m128 m = _mm_set_ps1(0.f);
    printf("Hello\n");
    ...
}

```

i686 Prologue:

```

    pushl   %edi
    leal    8(%esp), %edi
    andl    $-16, %esp
    pushl   4(%edi)
    pushl   %ebp
    movl    %esp, %ebp
    subl    $24, %esp
    call    .L1

```

```

.L1:
    popl    %ebx

```

i686 Body:

```

    movl    %edi, vreg1
    movl    (vreg1), %eax
    subl    %eax, %esp
    andl    $-16, %esp
    movl    %esp, %eax

```

Body:

```

    call    printf@PLT

```

i686 Epilogue:

```

    movl    %ebp, %esp
    popl    %ebp
    leal    -8(%edi), %esp
    popl    %edi
    ret

```

m128 and nested function and PIC

```

void foo()
{
    void bar(int arg1, int arg 2)
    {
        volatile __m128 m = _mm_set_ps1(0.f);
        ...
    }
    bar(1,2);
}

```

i686:

```

foo:
    ...
    movl    %ebp, %ecx
    call    bar@PLT
    ...

```



```

bar:
    pushl    %edi
    leal     8(%esp), %edi
    andl     $-16, %esp
    pushl    4(%edi)
    pushl    %ebp
    movl     %esp, %ebp
    subl     $24, %esp
    call     .L1
.L1:
    popl     %ebx

    movl     %edi, %reg1
    movl     (%reg1), %eax
    subl     %eax, %esp
    andl     $-16, %esp
    movl     %esp, %eax
    ...

    movl     %ebp, %esp
    popl     %ebp
    leal     -8(%edi), %esp
    popl     %edi
    ret

```

ml28, dynamic stack alloc and register parameter function call

```
static void bar(int arg1, int arg 2, int arg3)
```

```

{
    char * ptr=alloca(size);
    volatile __m128 m = _mm_set_ps1(0.f);
    ...
}

```

```

void foo()
{
    bar(1,2,3);
}

```

```

i686 foo:
    movl     $1, %eax
    movl     $2, %edx
    movl     $3, %ecx
    call     bar
    ...

bar:
    pushl    %edi
    leal     8(%esp), %edi
    andl     $-16, %esp
    pushl    $4(%edi)
    pushl    %ebp
    movl     %esp, %ebp
    subl     $24, %esp

    movl     %edi, %reg1
    movl     (%reg1), %eax
    subl     %eax, %esp
    andl     $-16, %esp
    movl     %esp, %eax
    ...

    movl     %ebp, %esp
    popl     %ebp
    leal     -8(%edi), %esp
    popl     %edi
    ret

```

Thanks - Joey

- **Follow-Ups:**

- [Re: A proposal to align GCC stack](#)

- *From:* Christian SchÄler

Index Nav: [\[Date Index\]](#) [\[Subject Index\]](#) [\[Author Index\]](#) [\[Thread Index\]](#)

Message Nav: [\[Date Prev\]](#) [\[Date Next\]](#) [\[Thread Prev\]](#) [\[Thread Next\]](#)